

Praktikum 4.

Anomaly Detection

(Fraud Detection of Credit Card)

Random Forest dan Logistic Regression

By: Rastri Prathivi, M.Kom.

```
[ ] 1 # import library
    2 import warnings
    3 warnings.filterwarnings('ignore')
    4
    5 import numpy as np
    6 import pandas as pd
    7 import seaborn as sns
    8 import matplotlib.pyplot as plt
    9 from imblearn.over_sampling import SMOTE
   10 from sklearn.model_selection import train_test_split, GridSearchCV
   11 from sklearn.preprocessing import StandardScaler
   12 from sklearn.ensemble import RandomForestClassifier
   13 from sklearn.linear_model import LogisticRegression
   14 from sklearn.ensemble import GradientBoostingClassifier
   15 from sklearn.metrics import accuracy_score, roc_auc_score, plot_roc_curve, confusion_matrix, plot_confusion_matrix
   16
```

UNTUK READ DATASET SILAKAN
DISESUAIKAN SENDIRI
CODINGNYA SEHINGGA DATA
DAPAT DIBACA

```
[ ] 1 # Read the dataset
    2 df = pd.read_csv("creditcard.csv")
    3 df.head()
```

Out[2]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.161603
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.187183
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.041744
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.276232
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

```
[ ] 1 # Printing quick information about the dataset
    2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Time    284807 non-null  float64
 1   V1      284807 non-null  float64
 2   V2      284807 non-null  float64
 3   V3      284807 non-null  float64
 4   V4      284807 non-null  float64
 5   V5      284807 non-null  float64
 6   V6      284807 non-null  float64
 7   V7      284807 non-null  float64
 8   V8      284807 non-null  float64
 9   V9      284807 non-null  float64
10  V10     284807 non-null  float64
```

```
[ ] 1 # Checking missing values in each column
     2 df.isnull().sum()
```

```
[ ] 1 # Identify duplicate values and mark all the duplicates as true
     2 # https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.duplicated.html
     3 df[df.duplicated(keep=False)]
```

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7
32	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.7
33	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.7
34	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.6
35	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.6
112	74.0	1.038370	0.127486	0.184456	1.109950	0.441699	0.945283	-0.
...
283485	171627.0	-1.457978	1.378203	0.811515	-0.603760	-0.711883	-0.471672	-0.
284190	172233.0	-2.667936	3.160505	-3.355984	1.007845	-0.377397	-0.109730	-0.
284191	172233.0	-2.667936	3.160505	-3.355984	1.007845	-0.377397	-0.109730	-0.
284192	172233.0	-2.691642	3.123168	-3.339407	1.017018	-0.293095	-0.167054	-0.
284193	172233.0	-2.691642	3.123168	-3.339407	1.017018	-0.293095	-0.167054	-0.

1854 rows × 31 columns

Out[4]:

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
```

```
[ ] 1 # https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop\_duplicates.html
    2 # drop data duplicated
    3 df = df.drop_duplicates(keep='first')
```

```
[ ] 1 # Check the distribution of the credit card fraud cases
    2 class_proportion = df['Class'].value_counts()
    3 class_proportion
```

Out[7]:

0 283253


1 473

Name: Class, dtype: int64

```

1 # Plotting a barchart to see the the distribution of the credit card fraud cases
2 plt.style.use('seaborn')
3 font1 = {'family': 'serif',
4         'fontstyle': 'italic',
5         'fontsize': 16,
6         'fontweight': 'bold',
7         'color': 'DarkBlue'}
8 font2 = {'weight': 'bold', 'size': 12}
9 font3 = {'weight': 'normal', 'size': 12}
10
11 fig, ax = plt.subplots(figsize=(12, 8))
12 sns.barplot(class_proportion.index, class_proportion.values, palette='Set2')
13 ax.set_title('Distribution of Credit Card Fraud Class', fontdict=font1)
14 ax.set_xlabel('Fraud class', fontdict=font2)
15 ax.set_xticklabels(ax.get_xticklabels(), rotation=0)
16 ax.set_ylabel('Number of transactions', fontdict=font2)
17 ax.set_yscale('log')
18 handles, labels = ax.get_legend_handles_labels()
19 ax.legend(handles, labels=['0: Legitimate', '1: Fraudulent'], prop= font3,
20         title = 'Fraud Class:', title_fontsize=14,
21         frameon=True, facecolor='white')
22 plt.show()

```



Feature Engineering and Data Modeling

```
[ ] 1 # Check the proportion of the fraud cases and identify the imbalance
     2 df['Class'].value_counts(normalize=True)
```


```
[ ] 1 # Arrange the dataset into features matrix and target vector
     2 # Drop the 'Time' variable as it does not that much help our analysis
     3 X = df.drop(columns=['Time', 'Class'])
     4 y = df['Class']
```

```
[ ] 1 # Make a SMOTE instance, then fit and apply it in one step
     2 # to create an oversampled version of our dataset.
     3
     4 sm = SMOTE(sampling_strategy='auto', random_state=3, k_neighbors=5)
     5 X_oversampled , y_oversampled = sm.fit_resample(X, y)
```

```
[ ] 1 # Summarize the fraud class distribution of the new SMOTE-transformed dataset
     2 unique_original, counts_original = np.unique(y, return_counts=True)
     3 unique_oversampled, counts_oversampled = np.unique(y_oversampled, return_counts=True)
     4
     5 print('Original fraud class distribution:', dict(zip(unique_original, counts_original)))
     6 print('New transformed fraud class distribution:', dict(zip(unique_oversampled, counts_oversampled)))
```



```
1 # Visualize the SMOTE-transformed target variable
2 plt.style.use('seaborn')
3 font1 = {'family': 'serif',
4         'fontstyle': 'italic',
5         'fontsize': 16,
6         'fontweight': 'bold',
7         'color': 'DarkBlue'}
8 font2 = {'weight': 'bold', 'size': 12}
9
10 fig, ax = plt.subplots(figsize=(12, 8))
11 sns.countplot(y_oversampled, palette='Set2', ax=ax)
12 ax.set_title('Distribution of Imbalanced Fraud Class Transformed by SMOTE', fontdict=font1)
13 ax.set_xlabel('Fraud class', fontdict=font2)
14 ax.set_xticklabels(['Legitimate', 'Fraudulent'])
15 ax.set_ylabel('Number of transactions', fontdict=font2)
16 plt.show()
```



Random Forest Classifier(RFC)

```
[ ] 1 # Separate the transformed features matrix and target vector into random train and test subsets
    2 X_train, X_test, y_train, y_test = train_test_split(X_oversampled, y_oversampled, random_state=3)
```

```
[ ] 1 # Instantiate and fit the model
    2 rfc = RandomForestClassifier(n_estimators=150)
    3 rfc.fit(X_train, y_train)
```

RandomForestClassifier(n_estimators=150)


```
[ ] 1 # Model Evalution -classification accuracy
    2 training_rfc_accuracy = rfc.score(X_train, y_train)
    3 testing_rfc_accuracy = rfc.score(X_test, y_test)
    4
    5 print("Training RFC Accuracy:", training_rfc_accuracy)
    6 print("Testing RFC Accuracy:", testing_rfc_accuracy )
```

Training RFC Accuracy: 1.0

Testing RFC Accuracy: 0.9998658447894823



```
1 # Plotting the confusion matrix
2 fig, ax = plt.subplots(figsize=(8, 8))
3 font1 = {'family': 'serif',
4         'fontstyle': 'italic',
5         'fontsize': 16,
6         'fontweight': 'bold',
7         'color': 'DarkRed'}
8 font2 = {'weight': 'bold', 'size': 12}
9
10 sns.heatmap(confusion_matrix(y_test, rfc.predict(X_test)),
11             cmap='Blues',
12             square=True,
13             annot=True,
14             fmt='d',
15             cbar_kws={'shrink': 0.8},
16             xticklabels=['Predicted 0s', 'Predicted 1s'],
17             yticklabels=['Actual 0s', 'Actual 1s'])
18 ax.set_title('RFC Confusion Matrix', fontdict=font1)
19 plt.show()
```



```
[ ] 1 # Model evaluation - Sensitivity, Specificity and Precision
2
3 TN, FP, FN, TP = confusion_matrix(y_test, rfc.predict(X_test)).flatten()
4 print("True Negatives:", TN)
5 print("False Positives:", FP)
6 print("False Negatives:", FN)
7 print("True Positives:", TP)
8
9 sensitivity = TP/(TP + FN)
10 specificity = TN/(TN + FP)
11 precision = TP/(TP + FP)
12 print("\nSensitivity:", sensitivity)
13 print("Specificity:", specificity)
14 print("Precision:", precision)
```

True Negatives: 70609

False Positives: 19

False Negatives: 0

True Positives: 70999

Sensitivity: 1.0

Specificity: 0.9997309848785184

Precision: 0.9997324621926835

```
[ ] 1 # Check the predicted probabilities for every observation in the test data subset
    2 # Note that the default classification threshold is 0.5
    3
    4 testing_probabilities= rfc.predict_proba(X_test)
    5 testing_probabilities
```

```
[ ] 1 # Convert the testing probabilities into a dataframe
    2 testing_probabilities_df = pd.DataFrame(testing_probabilities, columns=['1 - p(X_test)', 'p(X_test)'])
    3 testing_probabilities_df.head()
```

```
[ ] 1 # Get predictions
    2 rfc.predict(X_test)
```

Out[20]:

	1 - $p(X_{\text{test}})$	$p(X_{\text{test}})$
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	1.0	0.0
4	0.0	1.0

```
[ ] 1 # Model evaluation -AUC
    2 # Calculate AUC for both training and testing subsets
    3 # Only probabilities being in the positive class is needed for the calculation, that is the second column
    4 training_rfc_AUC = roc_auc_score(y_train, rfc.predict_proba(X_train)[:, 1])
    5 testing_rfc_AUC = roc_auc_score(y_test, rfc.predict_proba(X_test)[:, 1])
    6
    7 print("Training RFC AUC:", training_rfc_AUC)
    8 print("Testing RFC AUC:", testing_rfc_AUC)
```

Training RFC AUC: 1.0

Testing RFC AUC: 0.9999915933086132

Logist Regression (LGR) Model

```
[ ] 1 # Separate the transformed features matrix and target vector into random train and test subsets
     2 X_train, X_test, y_train, y_test = train_test_split(X_oversampled, y_oversampled, random_state=3)
```

```
[ ] 1 # define dictionary of hyperparameters
     2 params = {'penalty': ['l1', 'l2'],
             3         'C': [0.0001, 0.001, 0.01, 10, 50, 100],
             4         'class_weight': [None, 'balanced']}]
```

```
[ ] 1 # Instantiate Logistic Regression model. N.B: the default solver doesn't support l1 regularization
     2 # Instantiate Grid Search to find the best hyperparameters and fit the model
     3 lgr = LogisticRegression(solver='liblinear')
     4 gs = GridSearchCV(lgr, params, cv = 5)
     5 gs.fit(X_train, y_train)
```


```
[ ] 1 # Model evaluation - accuracy
     2 training_lgr_accuracy = gs.score(X_train, y_train)
     3 testing_lgr_accuracy = gs.score(X_test, y_test)
     4
     5 print("Training LGR Accuracy:", training_lgr_accuracy)
     6 print("Testing LGR Accuracy:", testing_lgr_accuracy)
```

Training LGR Accuracy: 0.958647050101323

Testing LGR Accuracy: 0.9577128654846887



```
1 # Plotting the confusion matrix
2 font1 = {'family': 'serif',
3          'fontstyle': 'italic',
4          'fontsize': 16,
5          'fontweight': 'bold',
6          'color': 'DarkRed'}
7
8 plot_confusion_matrix(gs, X_test, y_test, values_format='d')
9 plt.title('LGR Confusion Matrix', fontdict=font1)
10 plt.grid(False)
11 plt.show()
```



```
[ ] 1 # Model evaluation - Sensitivity, Specificity and Precision
    2 TN, FP, FN, TP = confusion_matrix(y_test, gs.predict(X_test)).flatten()
    3 print("True Negatives:", TN)
    4 print("False Positives:", FP)
    5 print("False Negatives:", FN)
    6 print("True Positives:", TP)
    7
    8 sensitivity = TP/(TP + FN)
    9 specificity = TN/(TN + FP)
   10 precision = TP/(TP + FP)
   11 print("\nSensitivity:", sensitivity)
   12 print("Specificity:", specificity)
   13 print("Precision:", precision)
```

True Negatives: 69349

False Positives: 1279

False Negatives: 4710

True Positives: 66289

Sensitivity: 0.9336610374794011

Specificity: 0.9818910347171094

Precision: 0.9810709211461046


```
[ ] 1 # Model evaluation -AUC
    2 # Calculate AUC for both training and testing subsets
    3 # Only probabilities being in the positive class is needed for the calculation, that is the second column
    4 training_lgr_AUC = roc_auc_score(y_train, gs.predict_proba(X_train)[:, 1])
    5 testing_lgr_AUC = roc_auc_score(y_test, gs.predict_proba(X_test)[:, 1])
    6
    7 print("Training LGR AUC:", training_lgr_AUC)
    8 print("Testing LGR AUC:", testing_lgr_AUC)
```

Training LGR AUC: 0.9918803796502905

Testing LGR AUC: 0.9915963855793377